

Incidents

- problems with our platform that impact its functionality, preventing customers from using the service they paid for (from a feature misbehaving to the full platform unavailability).
- they will happen, there's no way around it
- let's be ready for them and learn from them

Incident management

Infobip Engineering process, driven and owned by Team SRE

used to make sure production incidents are efficiently handled and provides a platform for learning by:

- making incidents transparent to all stakeholders
- driving resolution
- collecting data for reporting and insights
- providing a knowledge-base gained from incidents
- preventing incidents from repeating
- enabling discovery of systemic issues through long-term data tracking
- tracking and driving reliability improvements
- assuring compliance to laws, regulations and contracts

Incident Management stakeholders

- Customer Support
- Platform Operations
- Engineers
- Everyone else in Infobip

Phases of Incident Management

Identification

- make sure the problem is classified as an incident - does it impact the client?
- do the triage:
 - prove that the client is affected
 - which service is impacted
 - rough estimate of the impact
- report the incident ASAP
 - use #incident_management or IM Bot
- help customer support gather information so they can send notifications to clients

Intervention

- find and add engineers responsible for the affected service to the channel
- build a theory about the cause and gather data to support it
- try to fix the problem - keep everyone updated on what you are doing do the changes one by one
 - keep the changes documented
 - add other people if necessary

- assess the situation - is the incident likely to reoccur? do we have to monitor something new? do we have a plan if the incident happens again?

Review

- produce PIR
- detect contributing causes
- plan actions
- review the incident response as well as the incident itself